

I would deploy the Linear Regression model because it has the lowest test MAE (1.92) and essentially no overfitting, with a gap of -0.11. It is also simple, fast, and easy to maintain in production, which is an operational advantage over the more complex tree-based models. The trade-off is that it may miss nonlinear patterns that the Random Forest can capture, but the Random Forest's higher test MAE of 2.39 makes it not very suitable to deploy.